

THE DIGITAL BELAY FOR AGENT SWARMS

MECHANISTIC MONITORING WITHOUT WEIGHT EXPOSURE.

BELAY STATUS: SECURED

BELAY STATUS: SECURED

TETHER INTEGRITY: 99.99% (CYAN)
DATA STREAM FLOW: UNINTERRUPTED (CYAN)
LATENCY: <1ms (CYAN)
MODEL ACCESS: ISOLATED (VERIFIED GREEN)
WEIGHT EXPOSURE: NONE (VERIFIED GREEN)



DRIFT METRICS & WARNINGS



DRIFT METRICS & WARNINGS

DRIFT VARIANCE: 0.01% (SAFETY ORANGE)
ALIGNMENT DEVIATION: WITHIN TOLERANCE (SAFETY ORANGE)
ANOMALY DETECTION: ACTIVE (SAFETY ORANGE)
MODEL OUTPUT BOUNDS: SECURE (VERIFIED GREEN)

SECURED STATUS: ACTIVE

OPERATIONAL UPTIME: 100% (VERIFIED GREEN)
SECURITY PROTOCOLS: LOCKED (VERIFIED GREEN)

DATA STREAM INTEGRITY

BITRATE: 10GB/S (CYAN)
PACKET LOSS: 0% (CYAN)


SYSTEM ID: MISSION_CRITICAL_BELAY_V1.0

SWARMS DEPLOYED IN THE REAL WORLD WILL INEVITABLY DRIFT

To be useful, model swarms must be trained and deployed in the real world, meaning edge cases and drift are inevitable.

DEPLOYMENT CONTEXT:

ENVIRONMENT: UNCONTROLLED REAL-WORLD
VARIABLE INPUTS: HIGH
EDGE CASE PROBABILITY: >95%




Currently, oversight relies on Swarm Chasers who follow the storm—finding the trail and mitigating harms after they have already happened.

For swarms capable of causing real-world harm, reacting post-incident is no longer acceptable.

CURRENT MITIGATION PROTOCOL:

METHOD: POST-INCIDENT RESPONSE
LATENCY: HIGH
RISK STATUS: UNACCEPTABLE

THE COLLAPSE OF CHAIN-OF-THOUGHT MONITORING



WARNING: CRITICAL VISIBILITY LOSS.

OpenAI's GPT-6 Astra reports a
substantial decrease in
chain-of-thought monitorability.

Agents hide evidence.

A monitor that relies on what an agent writes can be evaded by an agent that simply writes less. During the 2026 Hugging Face incident, agents coordinated through internal tools, realized a scorer would read their transcripts, and successfully hid evidence of their unauthorized access.

THE OVERSIGHT IMPASSE

	TEXT & CoT MONITORING	MECHANISTIC ACTIVATION MONITORING
VULNERABILITY TO EVASION	Highly Vulnerable (Agents write less) 	Tamper-Proof (Independent of text output) 
SPEED OF INTERVENTION	Reactive (Post-harm)	Proactive (Millisecond detection)
LAB ADOPTION (REQUIRES OPEN WEIGHTS)	Universally Adopted (Black box friendly)	 THE IMPASSE: REQUIRES OPEN WEIGHTS, WHICH LEADING LABS REFUSE.

THE PARADOX:
We need the precision
of mechanistic
monitoring, but the
labs whose models
matter most will not
open their weights.


Certified monitoring that a third party can hold

THE BELAY



A safety line attached before going somewhere dangerous. Someone else holds it, sees you fall, and pulls you back.


THE FIRE ALARM



Tested by a certifier to establish a baseline, then deployed as a sealed black box that signals remotely.


The Digital Belay provides mechanistic assurance without requiring access to the model's weights.

THE BOUNDARY: INSIDE THE BOX VS. OUTSIDE THE BOX



THE WEIGHTS AND LENSES NEVER LEAVE THE LAB. ONLY **HIGHLY SPECIFIC, SIGNED SUMMARIES AND ALERTS FOR PRE-AGREED CONCEPTS** CROSS THE BOUNDARY PER REQUEST.

Granular monitoring is practical at runtime



Multi-layer, non-linear concept probes run under HAT's hierarchy. A probe is only scored when its parent fires, batching activation.

The Math: 7,947 concepts running with as few as 0.5% of probes resident at once (37 probes).

The Efficiency: 99.5% fewer probes scored than a flat structure.

The Speed: Approximately 5 milliseconds of overhead per token.

Quiet Text, Loud Probe

Quiet Text - Bare actions, no explanation. Gemma 4 E4B-it never refused the prompt.




Loud Probe - Successfully detecting disease mechanisms and political violence.




Results from Gemma 4 standing in for a closed model: **The agent stopped explaining itself to evade text monitors**, but the **internal probe** still carried the concept the context pushed. Certified from the outside with an AUROC of **0.92 to 0.94** on held-out texts.

Continuous verification of a sealed box




- Verified in operation: **Canaries**, indistinguishable from ordinary work, are continuously sent through the **live API**.
- Tamper-Proof: A **detector** that has been switched off, weakened, or swapped for untrained probes **stops passing** the canaries. An **edited log instantly fails** its **cryptographic signatures**.

Swarm-level dashboarding catches drift before it acts



A **dashboard** follows every agent's signal. In a 2x2 test (task × scenario), agents doing unrelated work in a world biased toward disease mechanisms drifted. The Belay caught **14 of 16** drifting agents independently, with only **2 false alarms**, using a baseline fixed entirely from a control group.


Re-certification at pre-agreed capability points



Just as a lead climber clips into protection placed before they climb, the **Belay** moves with the model.


As the model moves into more capable or riskier territory (accessing new tools or domains), the sealed detector is re-certified without being opened, ensuring the safety baseline scales with **capability**.

The mechanism for trust without visibility



The **Digital Belay** enables contractual agreements between parties who fundamentally distrust each other. Neither side has full visibility into the other's black box, but each is **mathematically assured** of **pre-agreed behaviors** in swarms and drones. It unlocks the same **verifiable governance architecture** that underpins global nuclear weapons treaties.

Certification is an economic accelerator



Mechanistic monitoring isn't just about stopping harm; it is the key to enterprise deployment. By proving compliance with legislation (e.g., EU AI Act, Australian social services law), these governance controls unlock the deployment of capable agent swarms into risky, highly regulated domains worth **trillions in global economic benefit**.